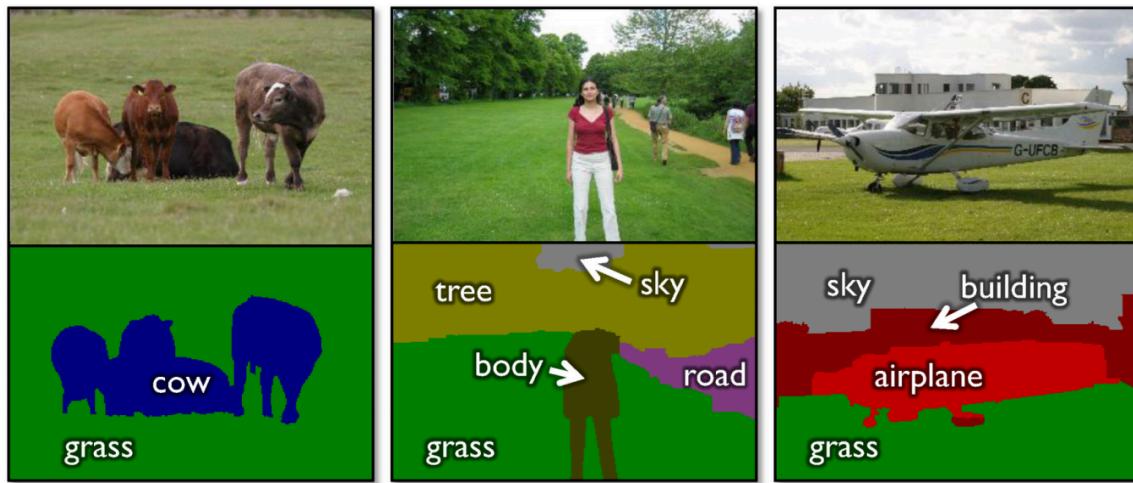


Semantic Segmentation, Dense Labeling

Liwei Wang

CS@UIUC

Semantic segmentation



| object classes | building | grass | tree | cow | sheep | sky | airplane | water | face | car |
|----------------|----------|-------|------|------|-------|------|----------|-------|------|------|
| bicycle | flower | sign | bird | book | chair | road | cat | dog | body | boat |

From [CS543 LAZ](#)

Why Semantic Segmentation ?

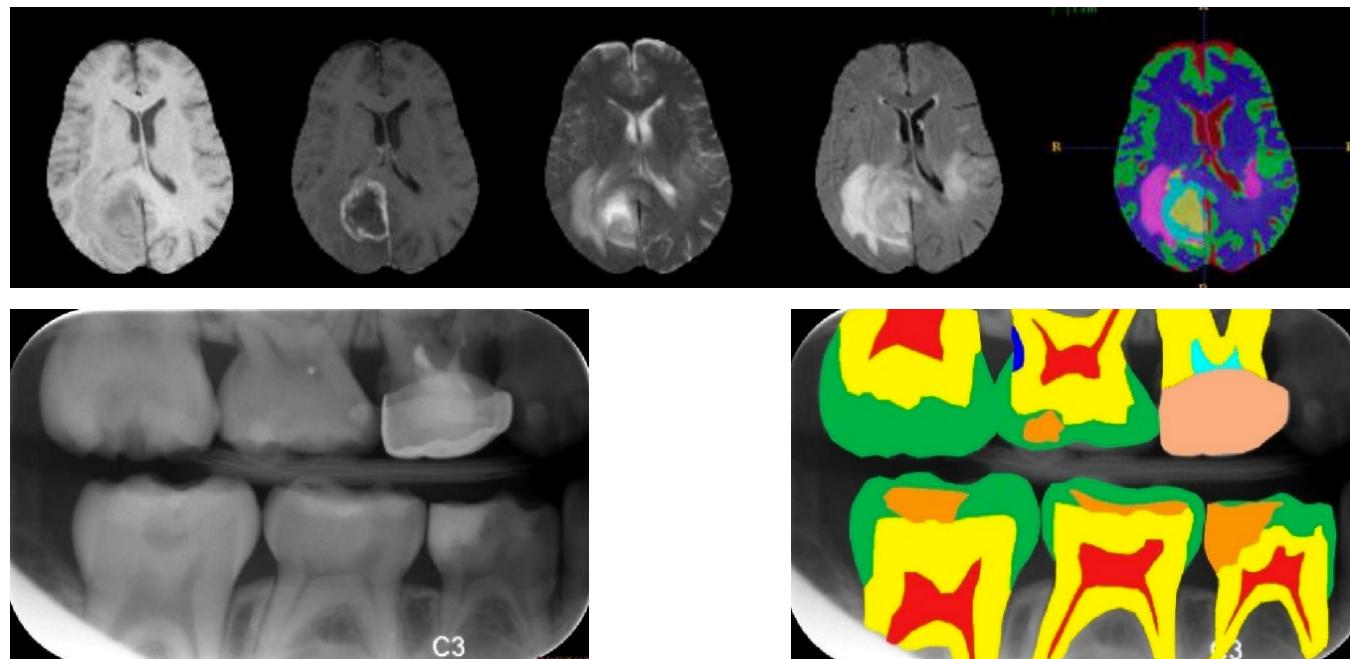
- Road Scene Understanding
- Useful for Self-Driving Car and autonomous drones



From [cityscape dataset](#)

Why Semantic Segmentation ?

- Medical Image Analysis



From [web](#)

Very Challenging Problem

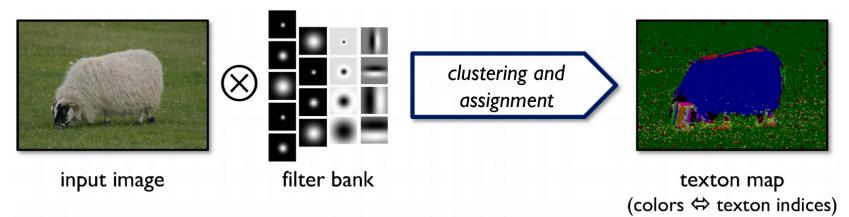


History

- Problem: label each pixel by one of C classes
- Define an energy function where unaries correspond to **local** classifier responses and smoothing potentials correspond to **contextual** terms
- Solve a multi-class graph cut problem

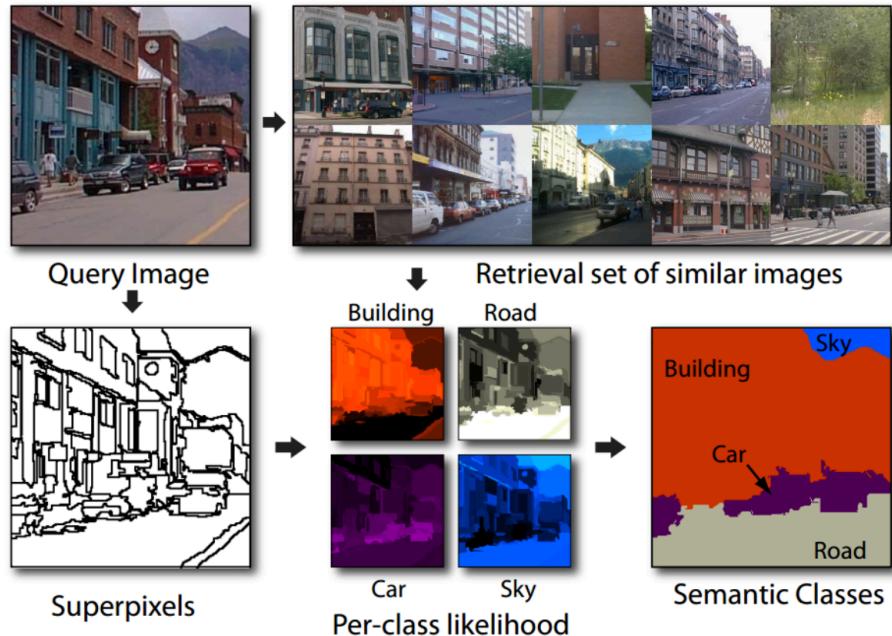
$$\log P(\mathbf{c}|\mathbf{x}, \theta) =$$

$$\sum_i \overbrace{\psi_i(c_i, \mathbf{x}; \theta_\psi)}^{\text{texture-layout}} + \overbrace{\pi(c_i, x_i; \theta_\pi)}^{\text{color}} + \overbrace{\lambda(c_i, i; \theta_\lambda)}^{\text{location}} \\ + \sum_{(i,j) \in \mathcal{E}} \overbrace{\phi(c_i, c_j, \mathbf{g}_{ij}(\mathbf{x}); \theta_\phi)}^{\text{edge}} - \log Z(\theta, \mathbf{x}) \quad (1)$$



From [TextonBoost, ECCV 2006](#)

History



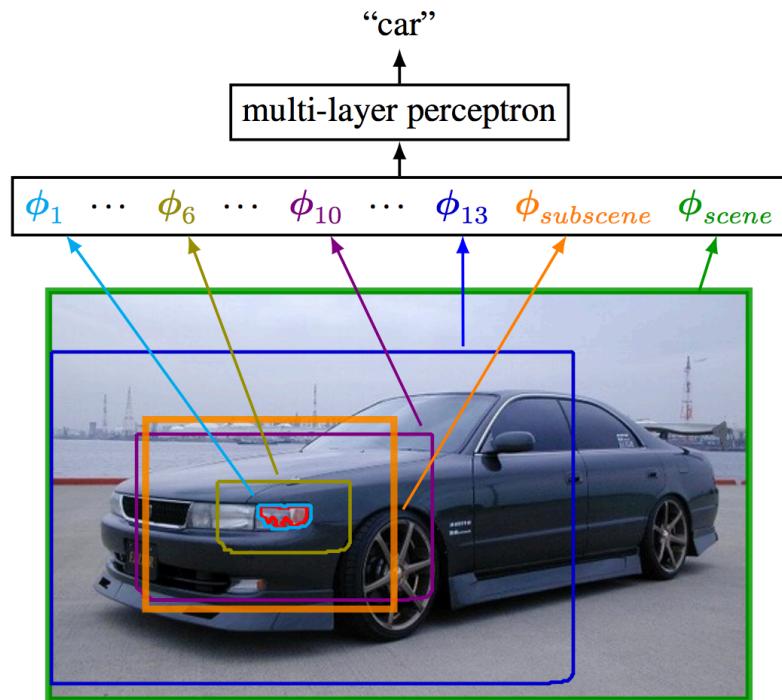
CRF energy function is defined on super-pixels:

- Unaries are based on nearest neighbor retrieval
- Pairwise potentials capture class co-occurrence statistics

Now, what is happening

- How deep neural networks can be used for Semantic Segmentation ?
- How to model local and contextual information with Deep Nets ?
- Differences and Similarities among methods ?

Zoom-out Features



- A simple **Feed-Forward Network**
- **Feature Concatenation** from Different Scales
- Strong features + Softmax Classification

Zoom-out Features

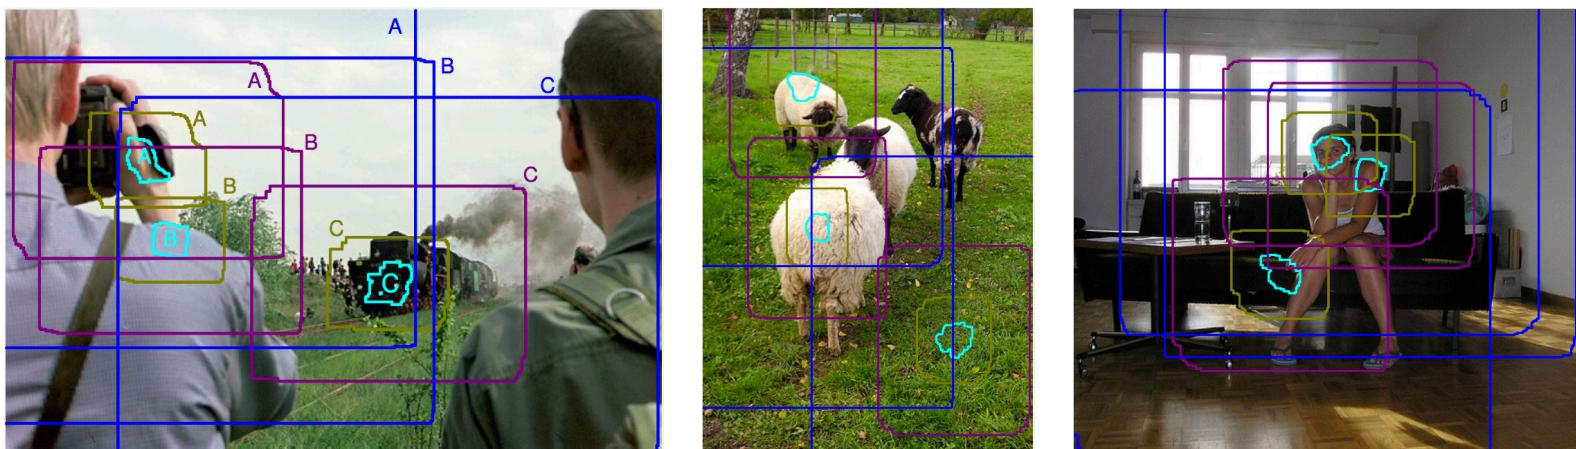
- **Sub-scene Level Features**

- Bounding box of superpixels within radius three from the superpixel at hand
- Warp bounding box to 256 x 256 pixels
- Activations of the last fully connected layer

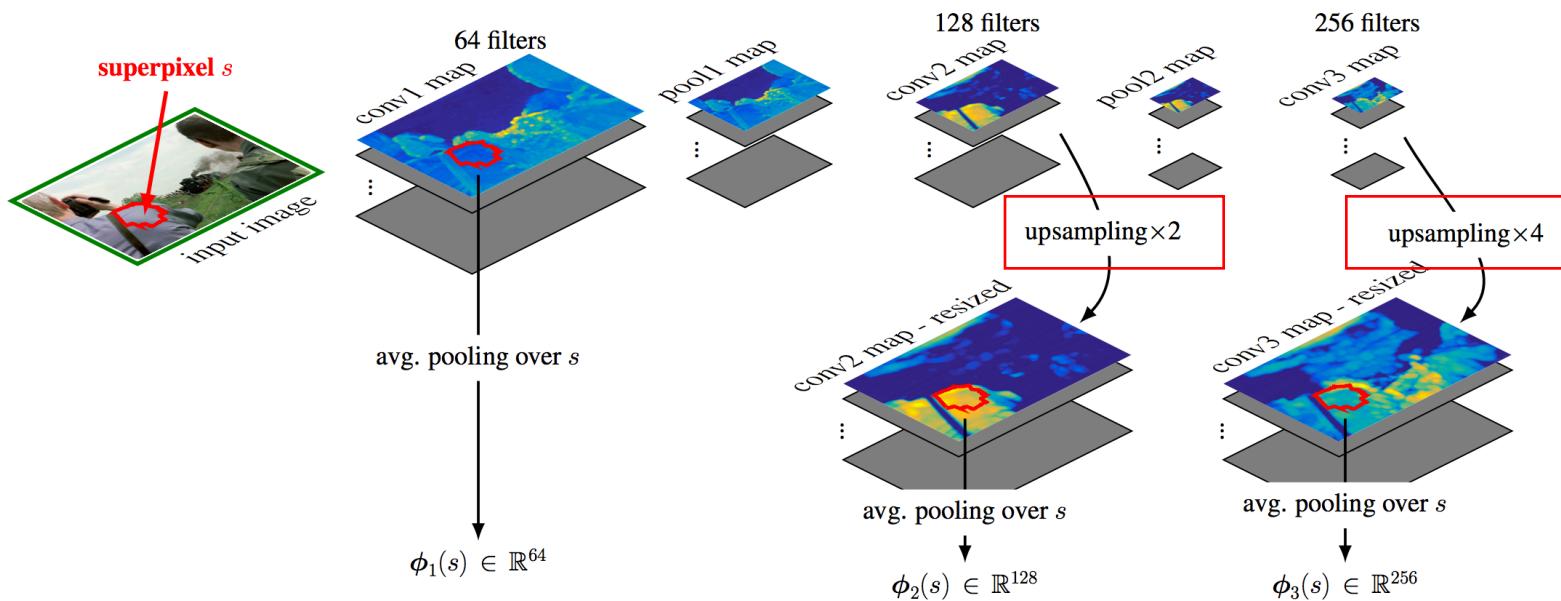
- **Scene Level Features**

- Warp image to 256 x 256 pixels
- Activations of the last fully connected layer

Zoom-out Features

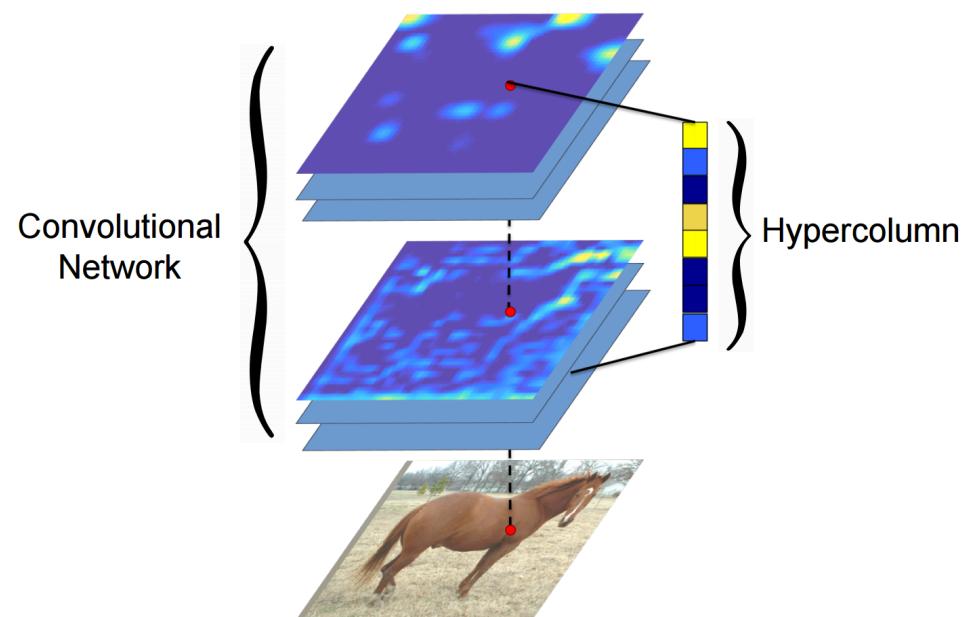


Zoom-out Features



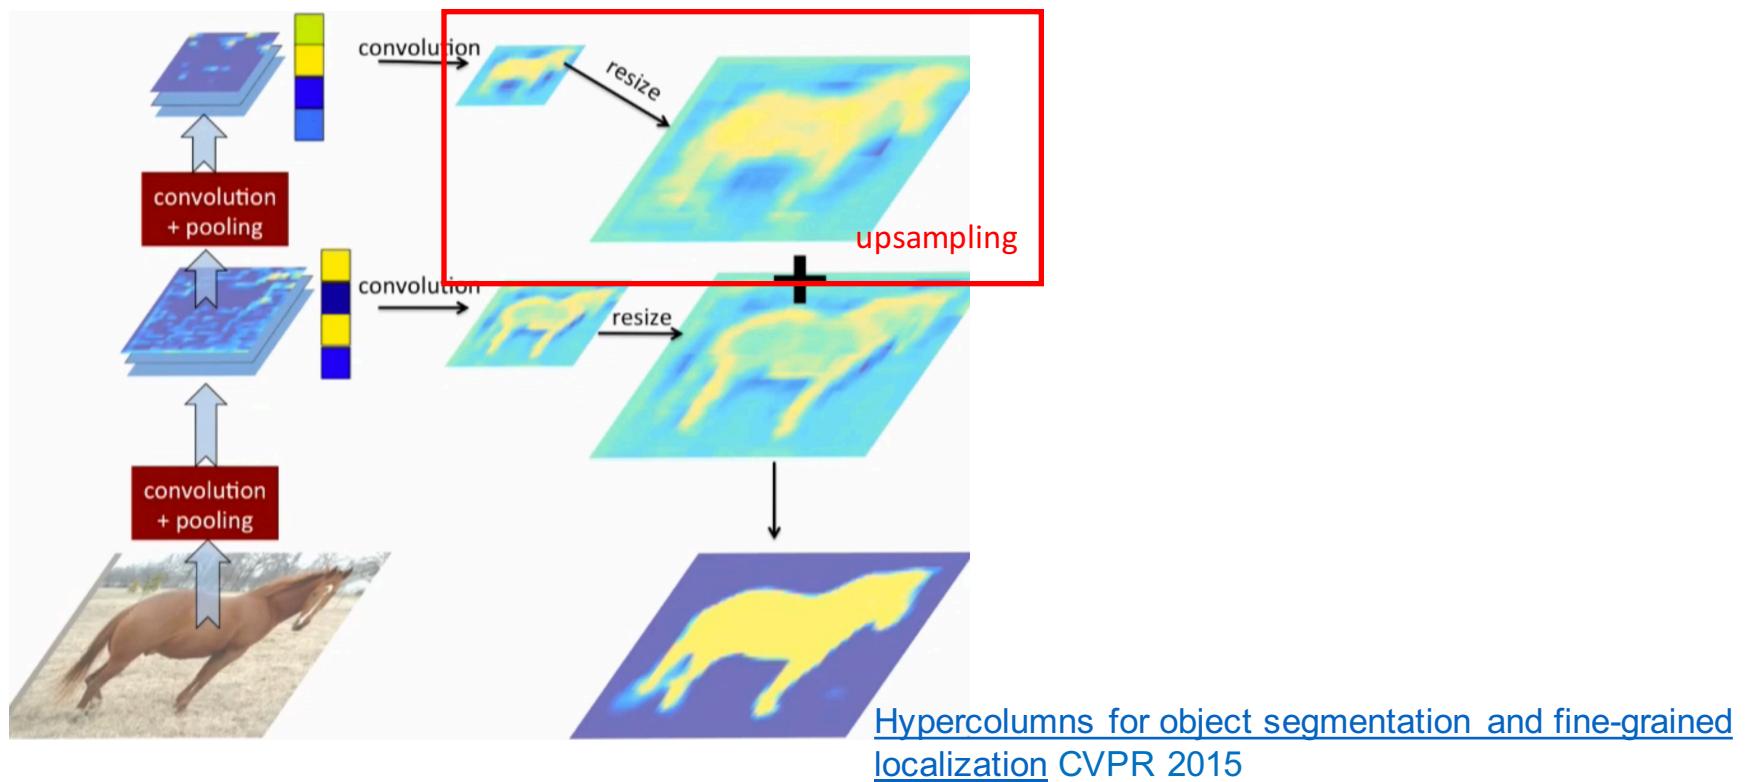
[Feedforward Semantic Segmentation With Zoom-Out Features, CVPR 2015](#)

Hypercolumns Representation

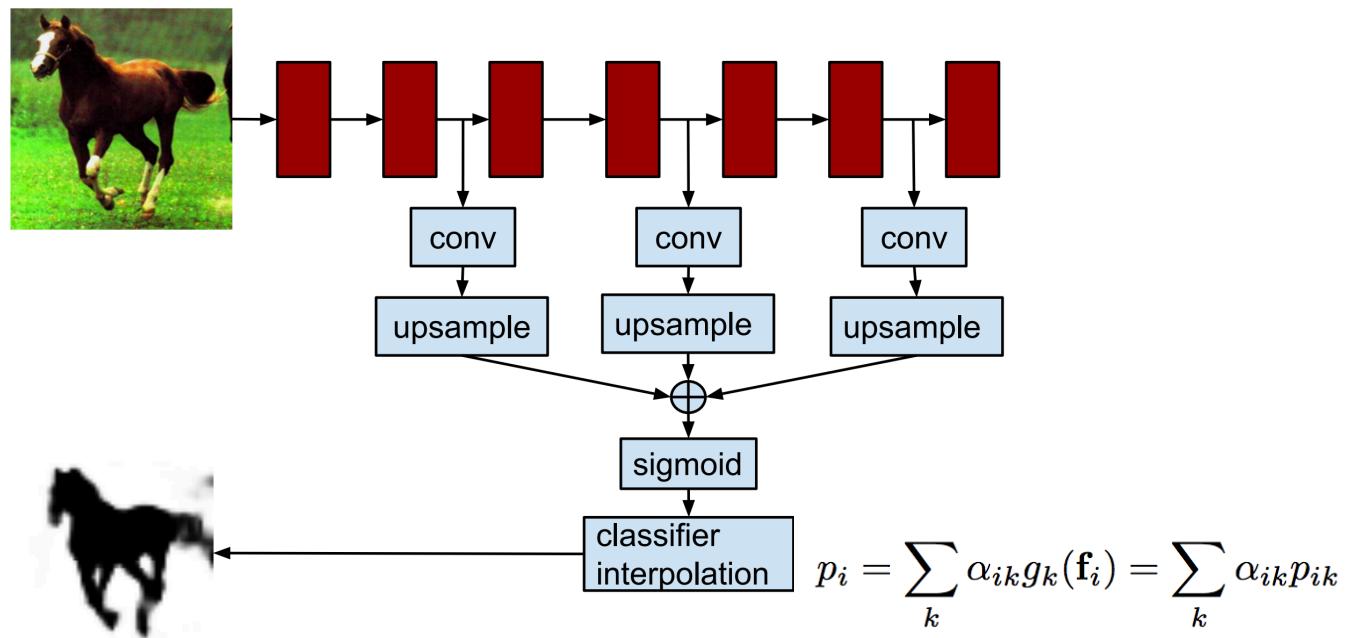


[Hypercolumns for object segmentation and fine-grained localization CVPR 2015](#)

Hypercolumn Representation



Hypercolumn Representation



[Hypercolumns for object segmentation and fine-grained localization CVPR 2015](#)

Evaluation

- Mean IoU

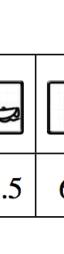
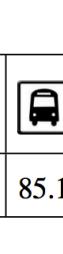
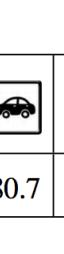
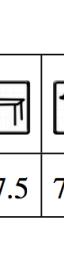
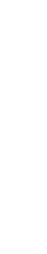
Per-class evaluation: an intersection of the predicted and true sets of pixels for a given class, divided by their union (IoU)

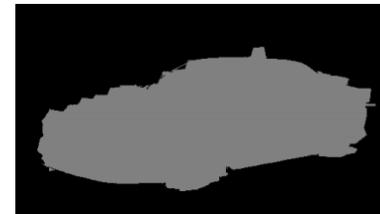
$$\text{seg. accuracy} = \frac{\text{true pos.}}{\text{true pos.} + \text{false pos.} + \text{false neg.}}$$

| | VOC 2012 |
|-------------|----------|
| Zoom-out | 69.6 |
| Hypercolumn | 62.6 |

Evaluation

Zoom-out Method on Pascal VOC2012

| class | mean | bg |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  | | | | | |
|-------|------|------|---|---|---|---|---|--|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| acc | 69.6 | 91.9 | 85.6 |  | 37.3 |  |  |  |  |  |  |  |  |  |  |  |  |  |  | | | | |

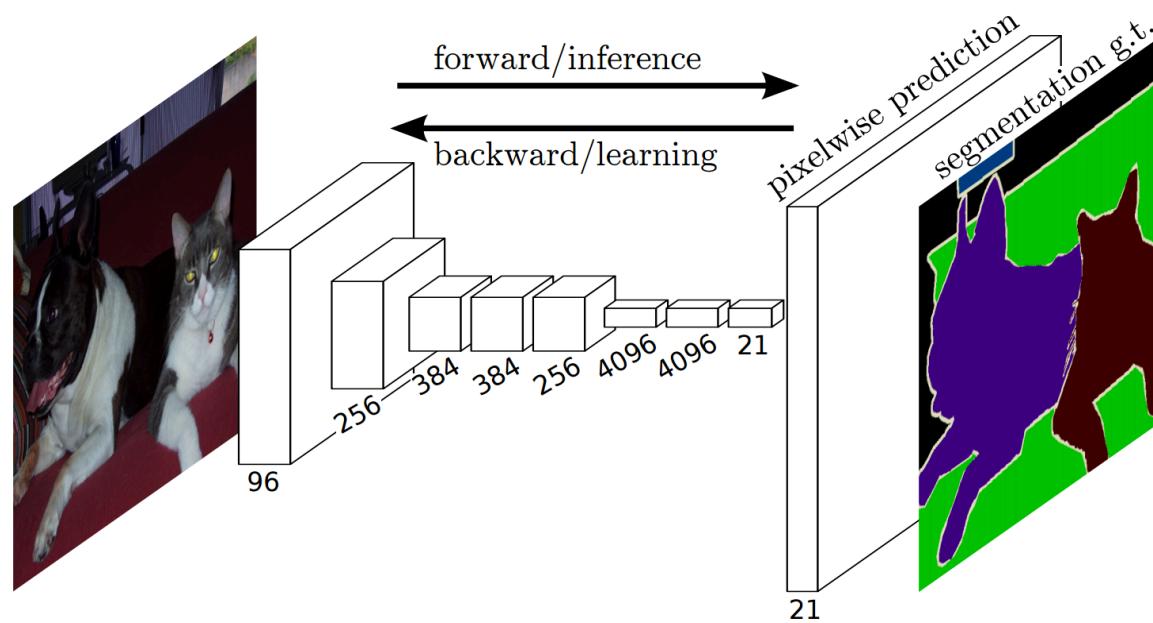


Hypercolumn and Zoom out

- Both uses the **multi-scale** features from intermediate layers in CNN
- Both use **upsampling** operations for each scale

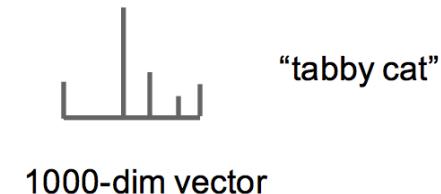
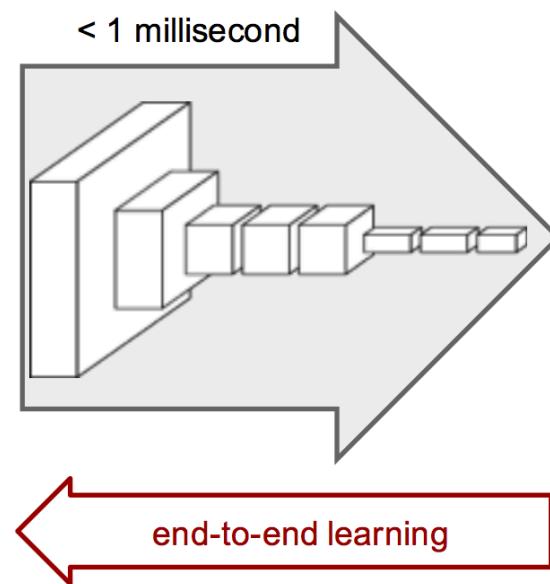
Any Pixel to Pixel ways ?
Can upsampling be learned ?

FCN for Semantic Segmentation



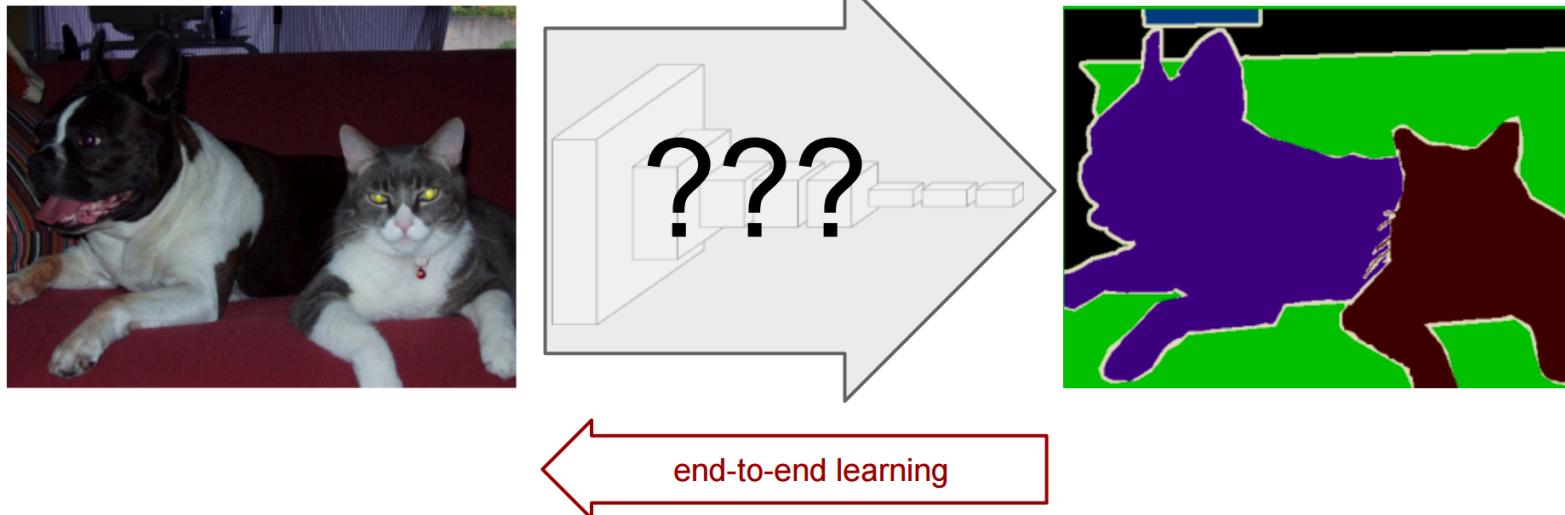
Fully convolutional Networks for Semantic Segmentation, CVPR 2015

Convnets for Classification



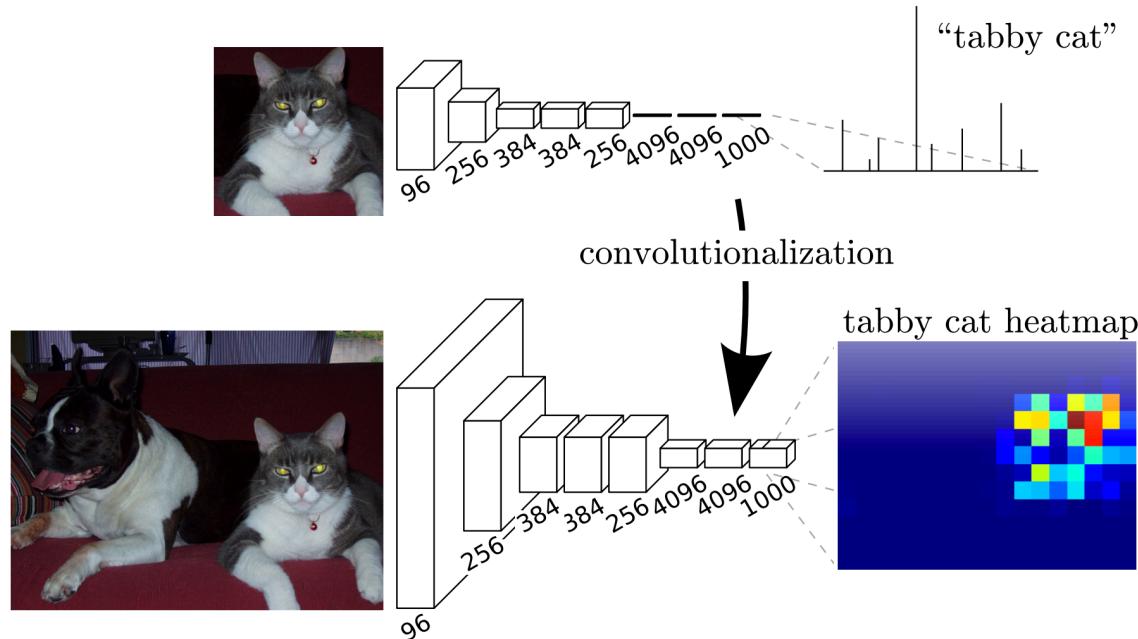
Fully convolutional Networks for Semantic Segmentation, CVPR 2015

Convnets for Segmentation ?



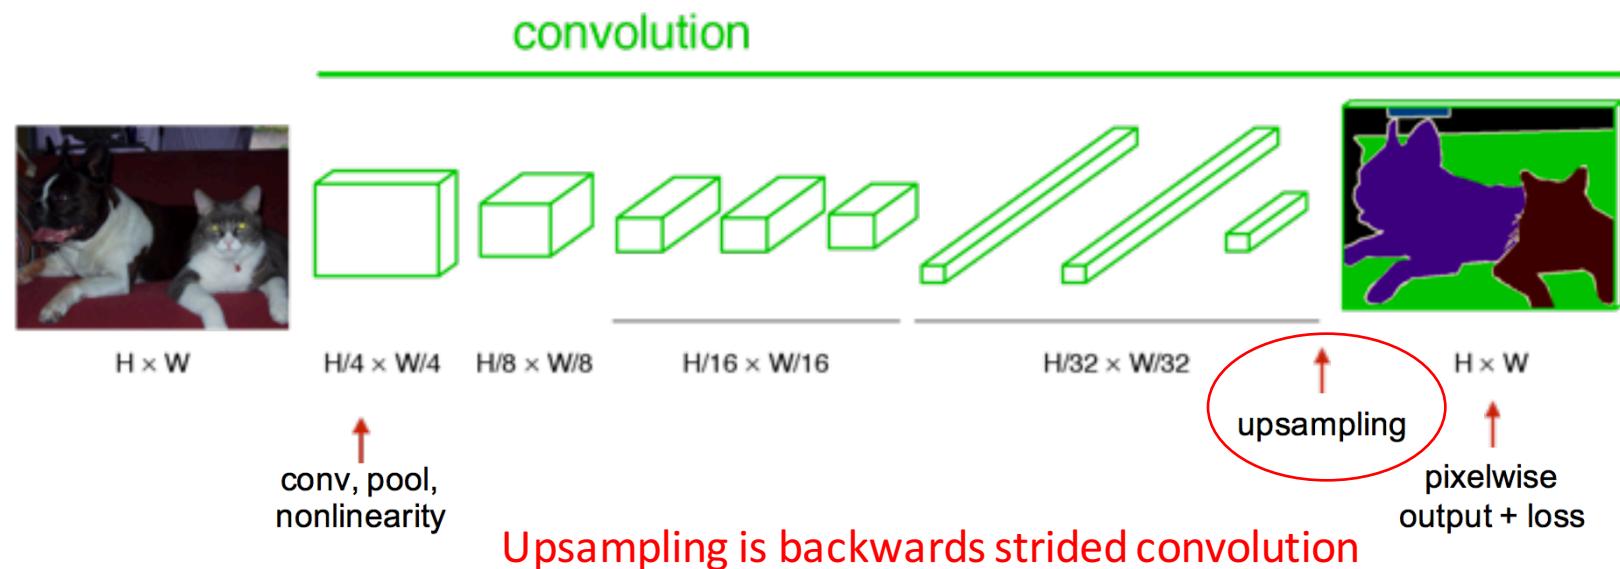
Fully convolutional Networks for Semantic Segmentation, CVPR 2015

From Convnets to FCN



Fully convolutional Networks for Semantic Segmentation, CVPR 2015

Pixel in , Pixel out



How does convolution works ?

Kernel/filter:

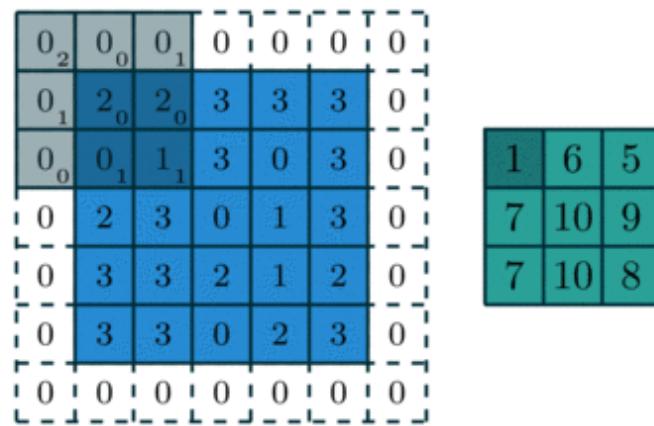
$$\begin{pmatrix} 0 & 1 & 2 \\ 2 & 2 & 0 \\ 0 & 1 & 2 \end{pmatrix}$$

| | | | | |
|----------------|----------------|----------------|---|---|
| 3 ₀ | 3 ₁ | 2 ₂ | 1 | 0 |
| 0 ₂ | 0 ₂ | 1 ₀ | 3 | 1 |
| 3 ₀ | 1 ₁ | 2 ₂ | 2 | 3 |
| 2 | 0 | 0 | 2 | 2 |
| 2 | 0 | 0 | 0 | 1 |

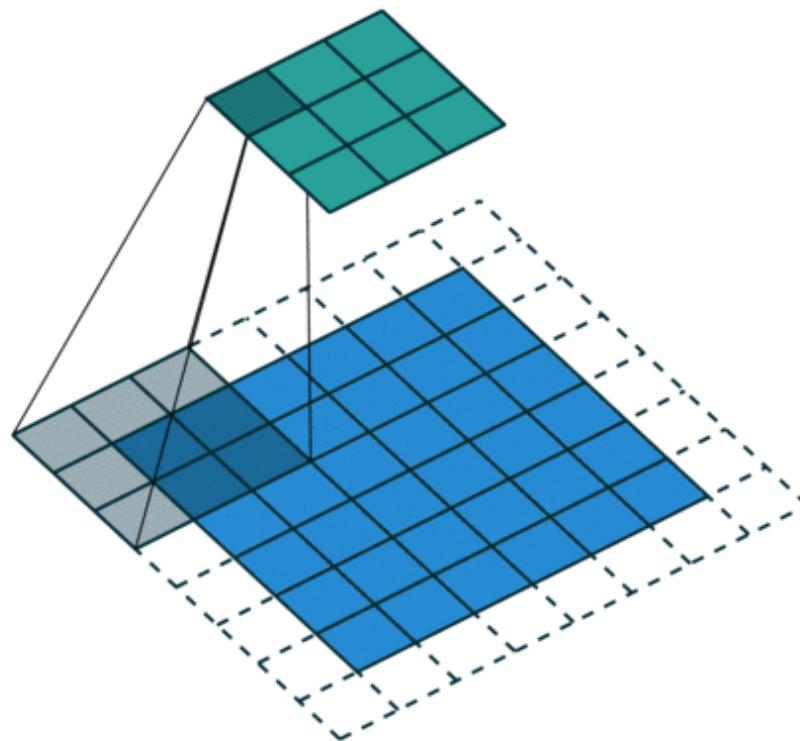
| | | |
|----|----|----|
| 12 | 12 | 17 |
| 10 | 17 | 19 |
| 9 | 6 | 14 |

How does convolution works ?

Kernel/filter:

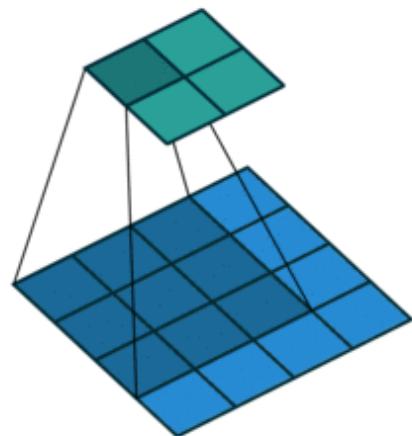


How does convolution works ?



How does convolution works ?

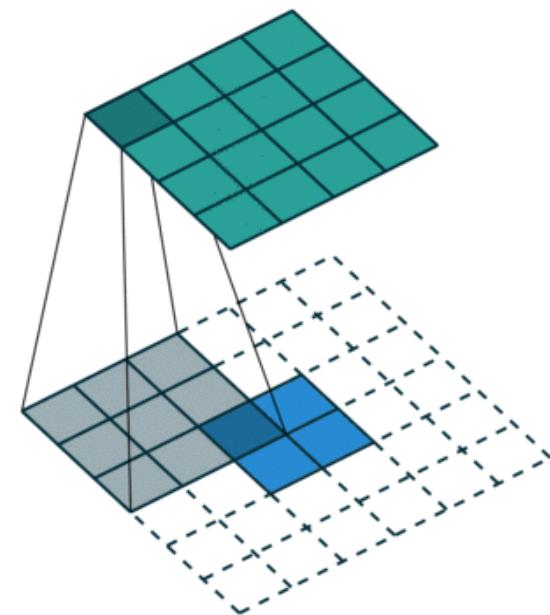
Convolution as a matrix operation



$$\begin{pmatrix} w_{0,0} & 0 & 0 & 0 \\ w_{0,1} & w_{0,0} & 0 & 0 \\ w_{0,2} & w_{0,1} & 0 & 0 \\ 0 & w_{0,2} & 0 & 0 \\ w_{1,0} & 0 & w_{0,0} & 0 \\ w_{1,1} & w_{1,0} & w_{0,1} & w_{0,0} \\ w_{1,2} & w_{1,1} & w_{0,2} & w_{0,1} \\ 0 & w_{1,2} & 0 & w_{0,2} \\ w_{2,0} & 0 & w_{1,0} & 0 \\ w_{2,1} & w_{2,0} & w_{1,1} & w_{1,0} \\ w_{2,2} & w_{2,1} & w_{1,2} & w_{1,1} \\ 0 & w_{2,2} & 0 & w_{1,2} \\ 0 & 0 & w_{2,0} & 0 \\ 0 & 0 & w_{2,1} & w_{2,0} \\ 0 & 0 & w_{2,2} & w_{2,1} \\ 0 & 0 & 0 & w_{2,2} \end{pmatrix}^T$$

From Theano Document Website

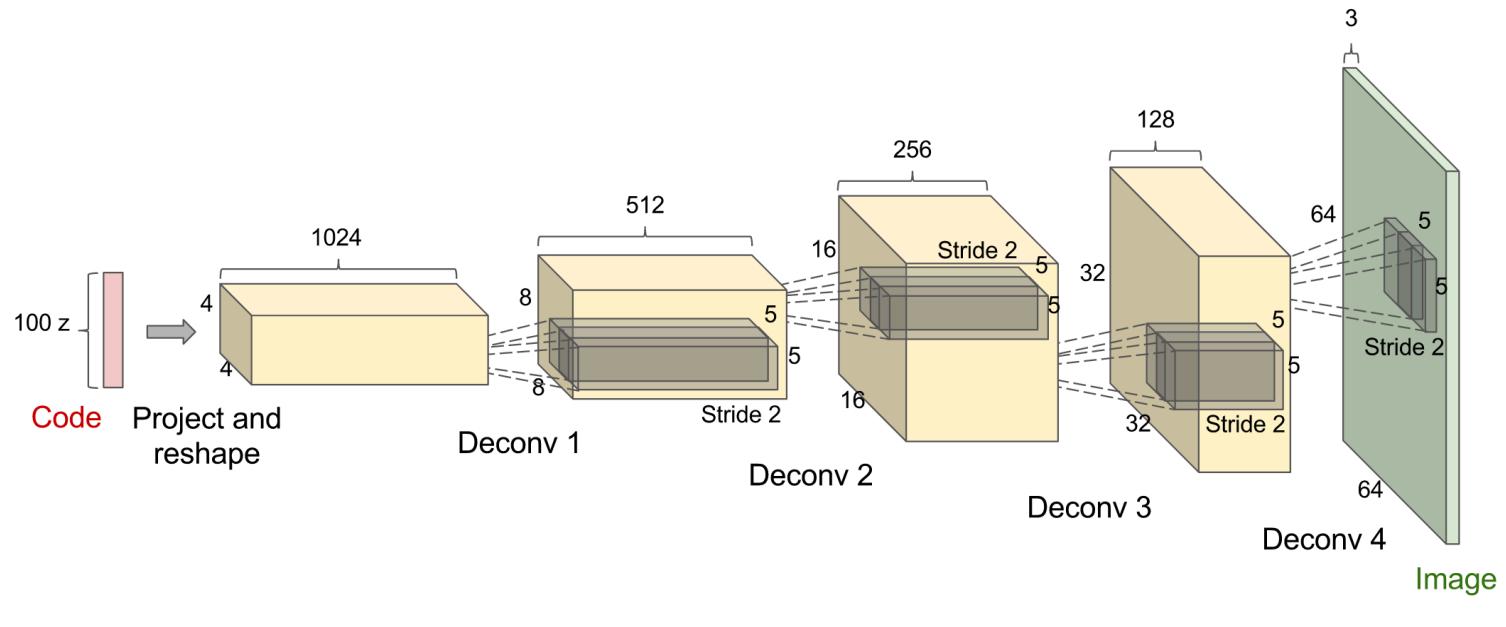
Upsampling is backwards strided convolution



Transposed convolution

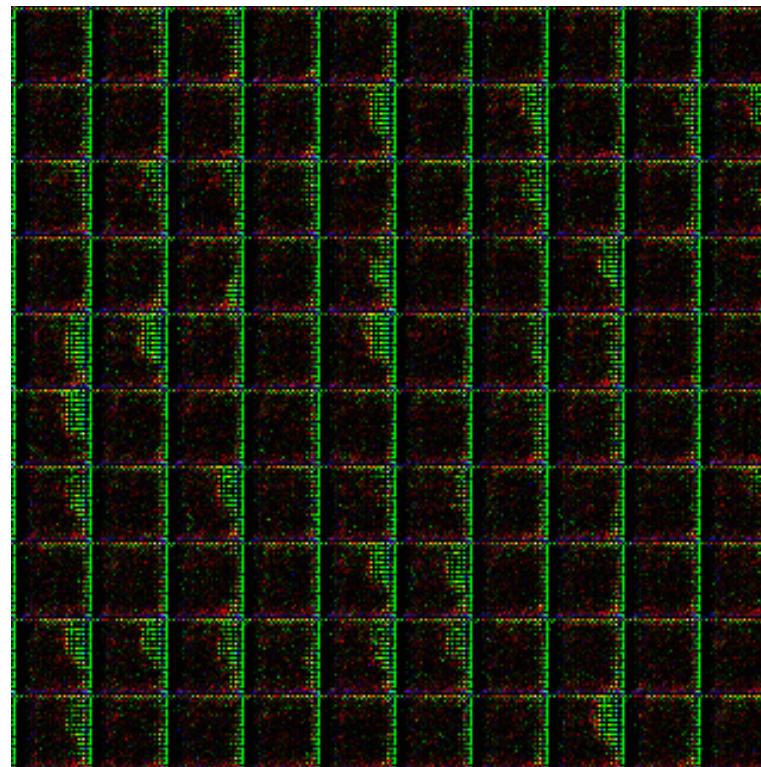
From Theano Document Website

Actually, we can build deconv networks...

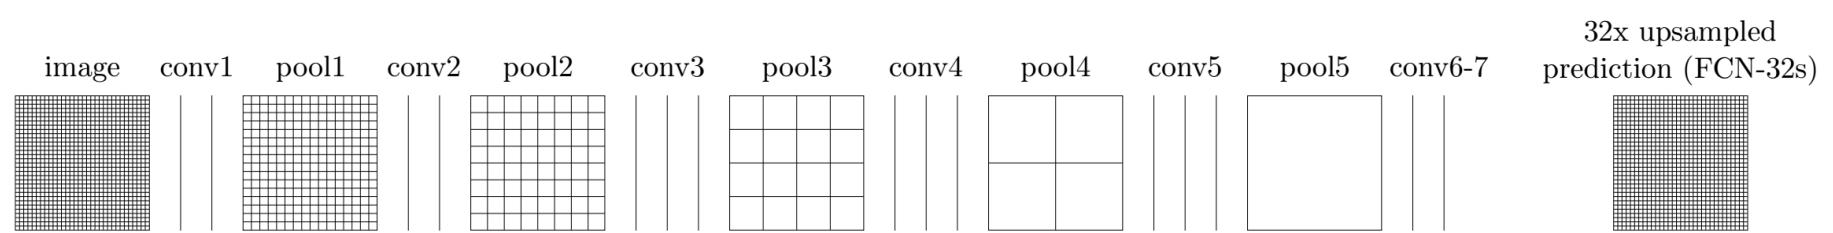


DCGAN

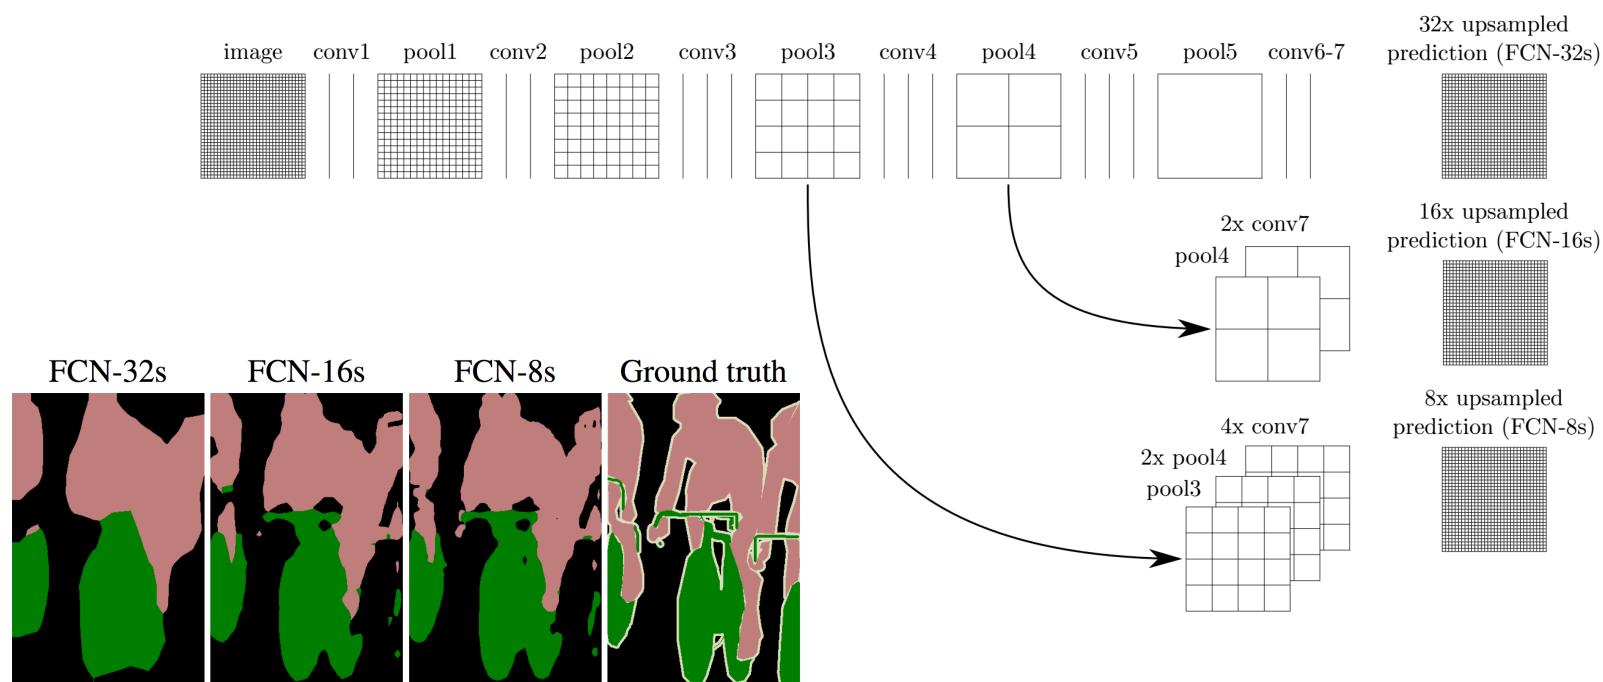
Deconv Layers for generating image !



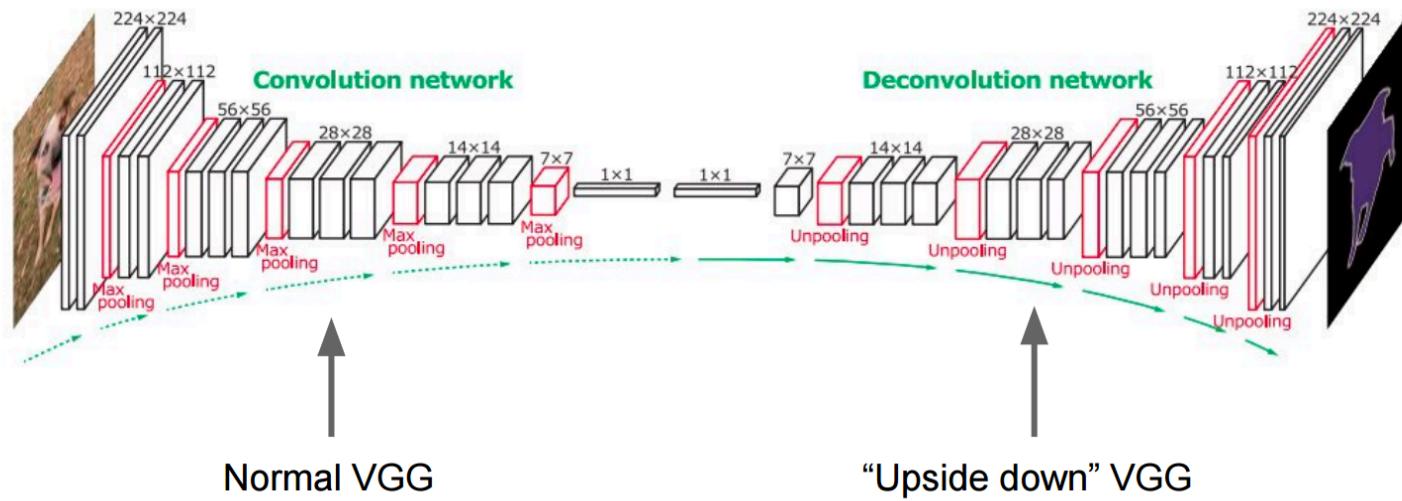
Framework of FCN



Framework of FCN



More than one upsampling layer



Noh et al, “Learning Deconvolution Network for Semantic Segmentation”, ICCV 2015

FCN is still not good ?

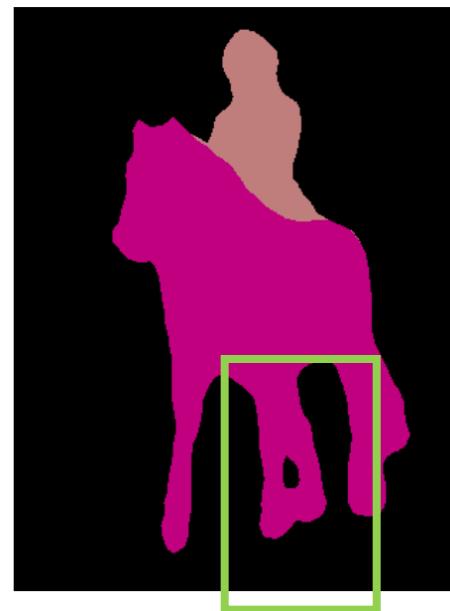
Ground Truth



Image



FCN-8s



Very coarse feature maps --- FCN-8s is still very coarse

Dilated Convolution

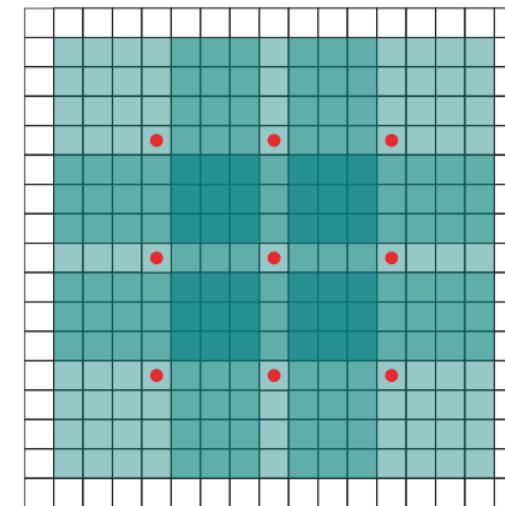
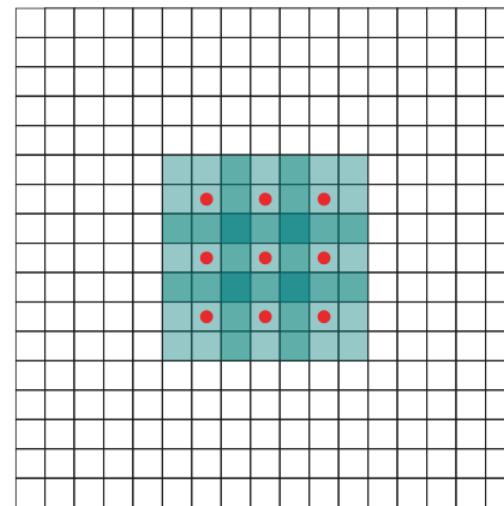
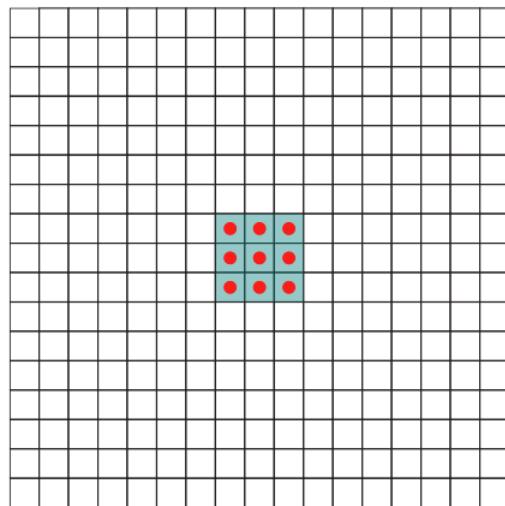
Discrete Convolution Operation:

$$(F * k)(\mathbf{p}) = \sum_{\mathbf{s} + \mathbf{t} = \mathbf{p}} F(\mathbf{s}) k(\mathbf{t})$$

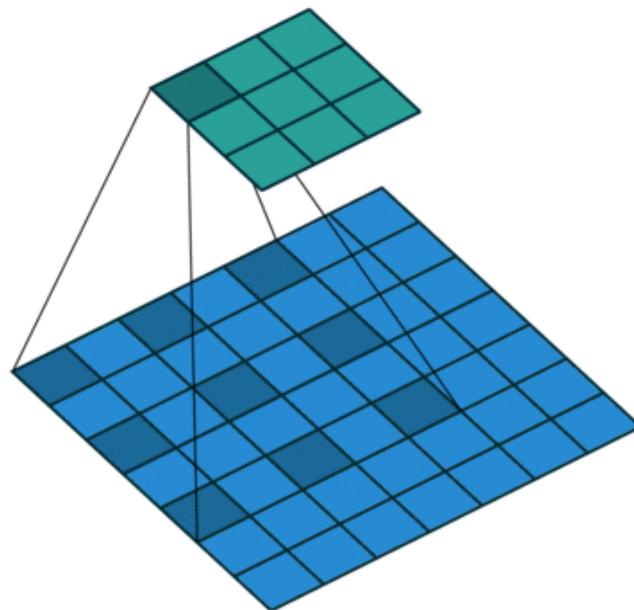
Dilated Convolution Operation:

$$(F *_l k)(\mathbf{p}) = \sum_{\mathbf{s} + l\mathbf{t} = \mathbf{p}} F(\mathbf{s}) k(\mathbf{t})$$

Dilated Convolution



Dilated Convolution



Network with Dilated Convolution

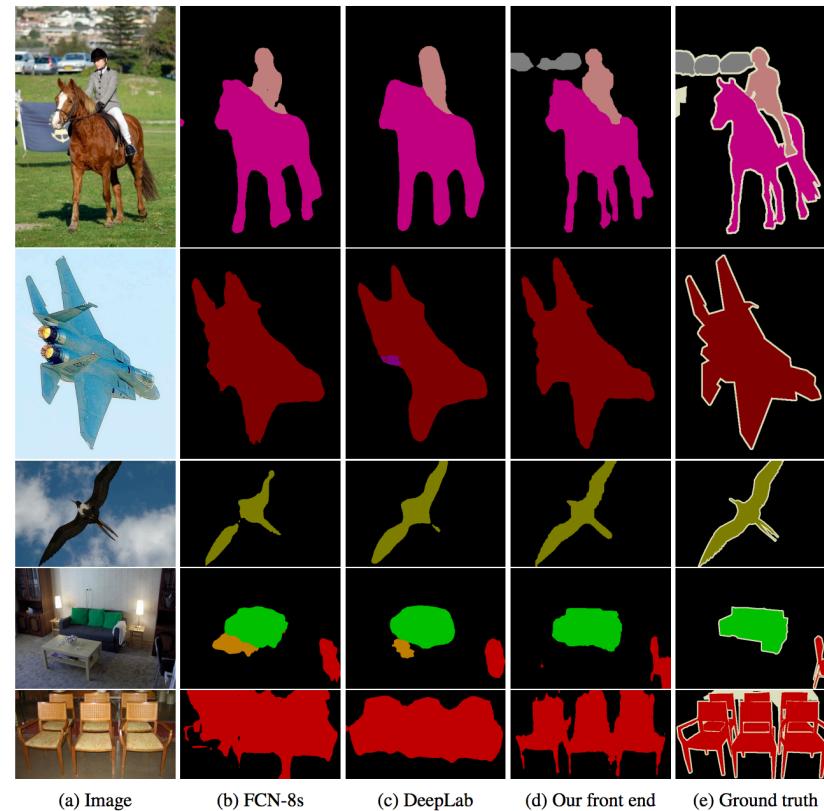
- Following FCN structure
- Using VGG-16 networks with modifications
- Called **Front End** and greatly improve performance

Network with Dilated Convolution

```
print 'VGG-16'
network = [
    {'k': 3, 'p': 1}, 'conv1_1',
    {'k': 3, 'p': 1}, 'conv1_2',
    {'k': 2, 's': 2}, 'pool1',
    {'k': 3, 'p': 1}, 'conv2_1',
    {'k': 3, 'p': 1}, 'conv2_2',
    {'k': 2, 's': 2}, 'pool2',
    {'k': 3, 'p': 1}, 'conv3_1',
    {'k': 3, 'p': 1}, 'conv3_2',
    {'k': 3, 'p': 1}, 'conv3_3',
    {'k': 2, 's': 2}, 'pool3',
    {'k': 3, 'p': 1}, 'conv4_1',
    {'k': 3, 'p': 1}, 'conv4_2',
    {'k': 3, 'p': 1}, 'conv4_3',
    {'k': 2, 's': 2}, 'pool4',
    {'k': 3, 'p': 1}, 'conv5_1',
    {'k': 3, 'p': 1}, 'conv5_2',
    {'k': 3, 'p': 1}, 'conv5_3',
    {'k': 2, 's': 2}, 'pool5',
    {'k': 7}, 'fc6',
    {'k': 1}, 'fc7',
    {'k': 1}, 'fc8',
]
print 'VGG-16 with Dilated Convs for Dense Prediction\n'
print 'Pascal VOC front end'
network = [
    {'k': 3}, 'conv1_1',
    {'k': 3}, 'conv1_2',
    {'k': 2, 's': 2}, 'pool1',
    {'k': 3}, 'conv2_1',
    {'k': 3}, 'conv2_2',
    {'k': 2, 's': 2}, 'pool2',
    {'k': 3}, 'conv3_1',
    {'k': 3}, 'conv3_2',
    {'k': 3}, 'conv3_3',
    {'k': 2, 's': 2}, 'pool3',
    {'k': 3}, 'conv4_1',
    {'k': 3}, 'conv4_2',
    {'k': 3}, 'conv4_3',
    {'k': 3, 'd': 2}, 'conv5_1',
    {'k': 3, 'd': 2}, 'conv5_2',
    {'k': 3, 'd': 2}, 'conv5_3',
    {'k': 7, 'd': 4}, 'fc6',
    {'k': 1}, 'fc7',
    {'k': 1}, 'fc-final',
]
]
```

From [arunmallya](#)

Network with Dilated Convolution



Network with Dilated Convolution

| | aero | bike | bird | boat | bottle | bus | car | cat | chair | cow | table | dog | horse | mbike | person | plant | sheep | sofa | train | tv | mean IoU |
|---------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-----------|-----------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| FCN-8s | 76.8 | 34.2 | 68.9 | 49.4 | 60.3 | 75.3 | 74.7 | 77.6 | 21.4 | 62.5 | 46.8 | 71.8 | 63.9 | 76.5 | 73.9 | 45.2 | 72.4 | 37.4 | 70.9 | 55.1 | 62.2 |
| DeepLab | 72 | 31 | 71.2 | 53.7 | 60.5 | 77 | 71.9 | 73.1 | 25.2 | 62.6 | 49.1 | 68.7 | 63.3 | 73.9 | 73.6 | 50.8 | 72.3 | 42.1 | 67.9 | 52.6 | 62.1 |
| DeepLab-Msc | 74.9 | 34.1 | 72.6 | 52.9 | 61.0 | 77.9 | 73.0 | 73.7 | 26.4 | 62.2 | 49.3 | 68.4 | 64.1 | 74.0 | 75.0 | 51.7 | 72.7 | 42.5 | 67.2 | 55.7 | 62.9 |
| Our front end | 82.2 | 37.4 | 72.7 | 57.1 | 62.7 | 82.8 | 77.8 | 78.9 | 28 | 70 | 51.6 | 73.1 | 72.8 | 81.5 | 79.1 | 56.6 | 77.1 | 49.9 | 75.3 | 60.9 | 67.6 |

Multi-scale Context Aggregation

| Layer | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
|-----------------|--------------|--------------|--------------|----------------|----------------|----------------|----------------|----------------|
| Convolution | 3×3 | 3×3 | 3×3 | 3×3 | 3×3 | 3×3 | 3×3 | 1×1 |
| Dilation | 1 | 1 | 2 | 4 | 8 | 16 | 1 | 1 |
| Truncation | Yes | Yes | Yes | Yes | Yes | Yes | Yes | No |
| Receptive field | 3×3 | 5×5 | 9×9 | 17×17 | 33×33 | 65×65 | 67×67 | 67×67 |
| Output channels | | | | | | | | |
| Basic | C | C | C | C | C | C | C | C |
| Large | $2C$ | $2C$ | $4C$ | $8C$ | $16C$ | $32C$ | $32C$ | C |

Multi-scale Context Aggregation

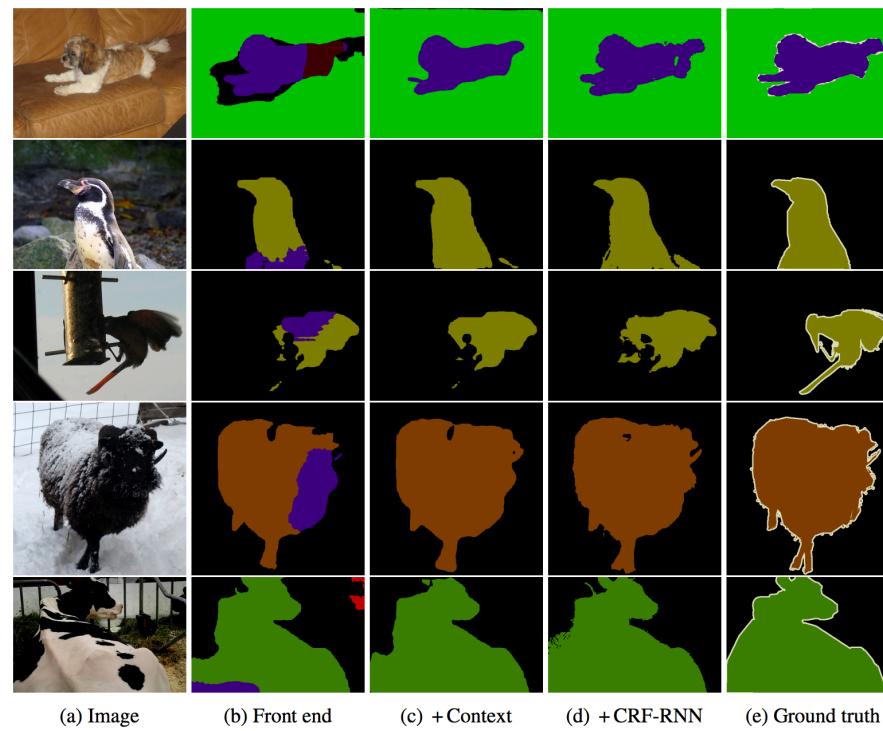
| Context Aggregation Module | | | |
|----------------------------|-----------------|------------------|-------------|
| Layer Name | Receptive Field | Effective Stride | Output Size |
| Input | -- | -- | 66 |
| ct_conv1_1 | 3 | 1 | 130 |
| ct_conv1_2 | 5 | 1 | 128 |
| ct_conv2_1 | 9 | 1 | 124 |
| ct_conv3_1 | 17 | 1 | 116 |
| ct_conv4_1 | 33 | 1 | 100 |
| ct_conv5_1 | 65 | 1 | 68 |
| ct_fc1 | 67 | 1 | 66 |
| ct_final | 67 | 1 | 66 |

From [arunmallya](#)

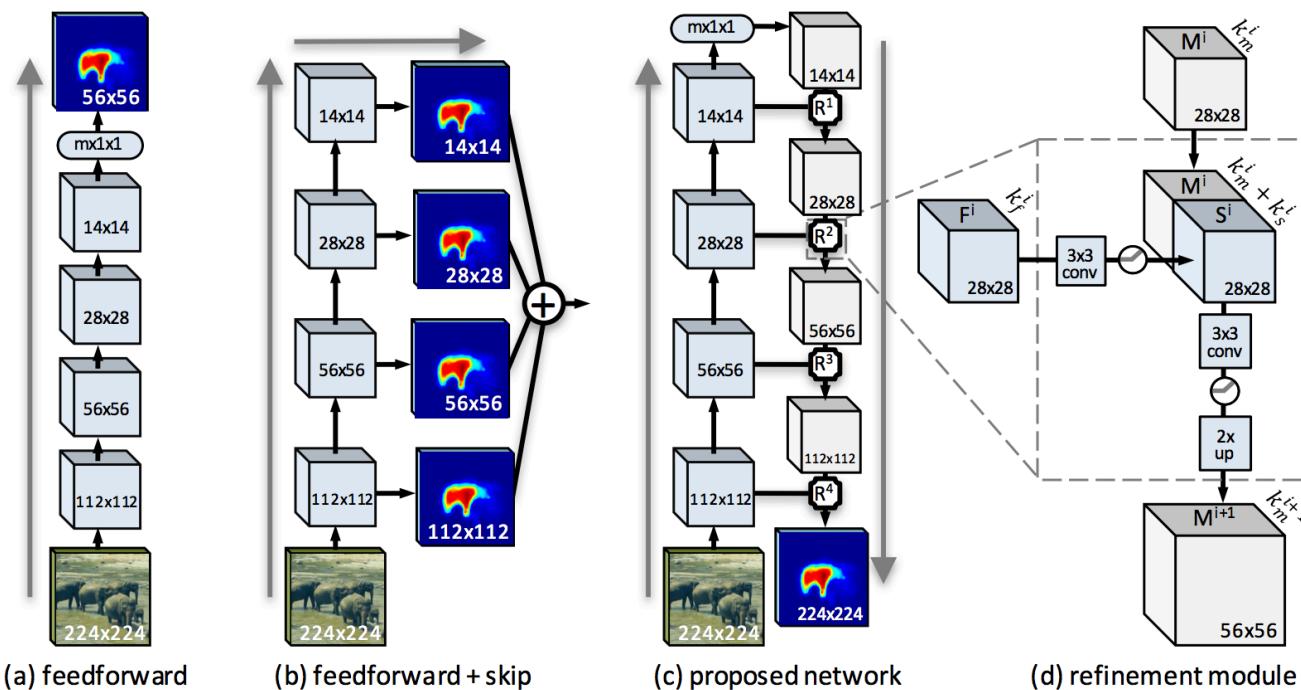
Front End + Context Net

| | aero | bike | bird | boat | bottle | bus | car | cat | chair | cow | table | dog | horse | mbike | person | plant | sheep | sofa | train | tv | mean IoU |
|---------------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| Front end | 86.3 | 38.2 | 76.8 | 66.8 | 63.2 | 87.3 | 78.7 | 82 | 33.7 | 76.7 | 53.5 | 73.7 | 76 | 76.6 | 83 | 51.9 | 77.8 | 44 | 79.9 | 66.3 | 69.8 |
| Front + Basic | 86.4 | 37.6 | 78.5 | 66.3 | 64.1 | 89.9 | 79.9 | 84.9 | 36.1 | 79.4 | 55.8 | 77.6 | 81.6 | 79 | 83.1 | 51.2 | 81.3 | 43.7 | 82.3 | 65.7 | 71.3 |
| Front + Large | 87.3 | 39.2 | 80.3 | 65.6 | 66.4 | 90.2 | 82.6 | 85.8 | 34.8 | 81.9 | 51.7 | 79 | 84.1 | 80.9 | 83.2 | 51.2 | 83.2 | 44.7 | 83.4 | 65.6 | 72.1 |
| Front end + CRF | 89.2 | 38.8 | 80 | 69.8 | 63.2 | 88.8 | 80 | 85.2 | 33.8 | 80.6 | 55.5 | 77.1 | 80.8 | 77.3 | 84.3 | 53.1 | 80.4 | 45 | 80.7 | 67.9 | 71.6 |
| Front + Basic + CRF | 89.1 | 38.7 | 81.4 | 67.4 | 65 | 91 | 81 | 86.7 | 37.5 | 81 | 57 | 79.6 | 83.6 | 79.9 | 84.6 | 52.7 | 83.3 | 44.3 | 82.6 | 67.2 | 72.7 |
| Front + Large + CRF | 89.6 | 39.9 | 82.7 | 66.7 | 67.5 | 91.1 | 83.3 | 87.4 | 36 | 83.3 | 52.5 | 80.7 | 85.7 | 81.8 | 84.4 | 52.6 | 84.4 | 45.3 | 83.7 | 66.7 | 73.3 |
| Front end + RNN | 88.8 | 38.1 | 80.8 | 69.1 | 65.6 | 89.9 | 79.6 | 85.7 | 36.3 | 83.6 | 57.3 | 77.9 | 83.2 | 77 | 84.6 | 54.7 | 82.1 | 46.9 | 80.9 | 66.7 | 72.5 |
| Front + Basic + RNN | 89 | 38.4 | 82.3 | 67.9 | 65.2 | 91.5 | 80.4 | 87.2 | 38.4 | 82.1 | 57.7 | 79.9 | 85 | 79.6 | 84.5 | 53.5 | 84 | 45 | 82.8 | 66.2 | 73.1 |
| Front + Large + RNN | 89.3 | 39.2 | 83.6 | 67.2 | 69 | 92.1 | 83.1 | 88 | 38.4 | 84.8 | 55.3 | 81.2 | 86.7 | 81.3 | 84.3 | 53.6 | 84.4 | 45.8 | 83.8 | 67 | 73.9 |

Front End + Context Net



Top to Down Refinement



Learning to refine object segments, ECCV 2016

Refine Deep Mask

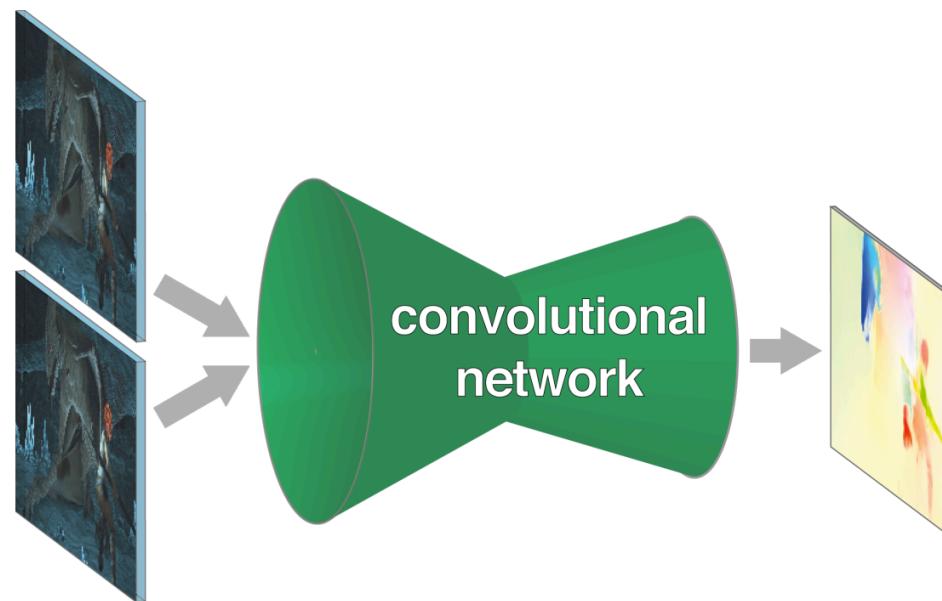


(a) DeepMask Output

(b) SharpMask Output

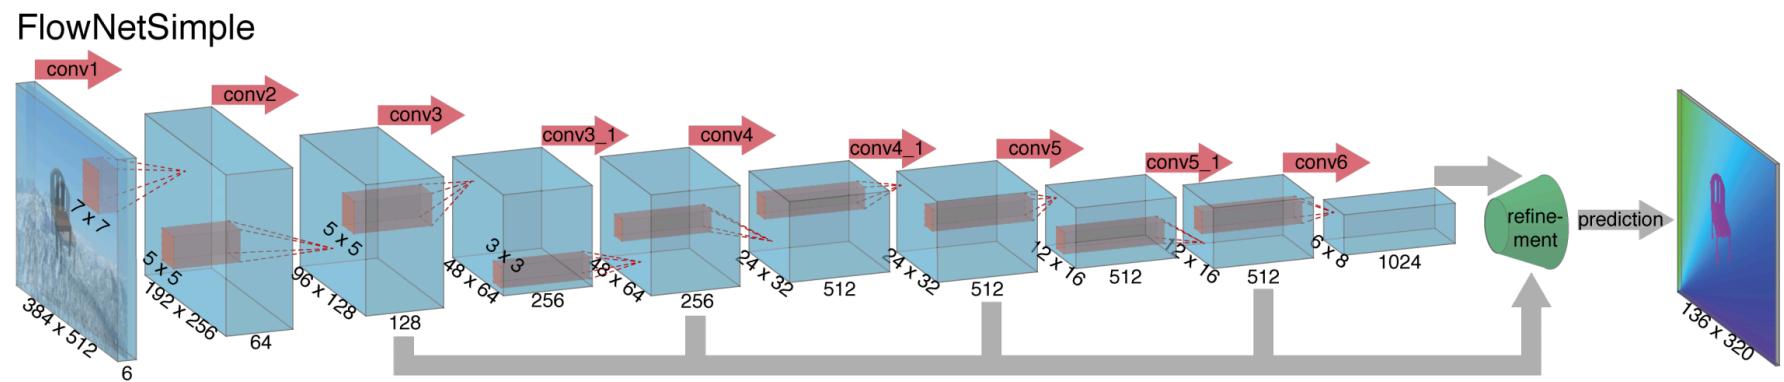
Learning to refine object segments, ECCV 2016

Dense Labeling Task: Learning Optical Flow



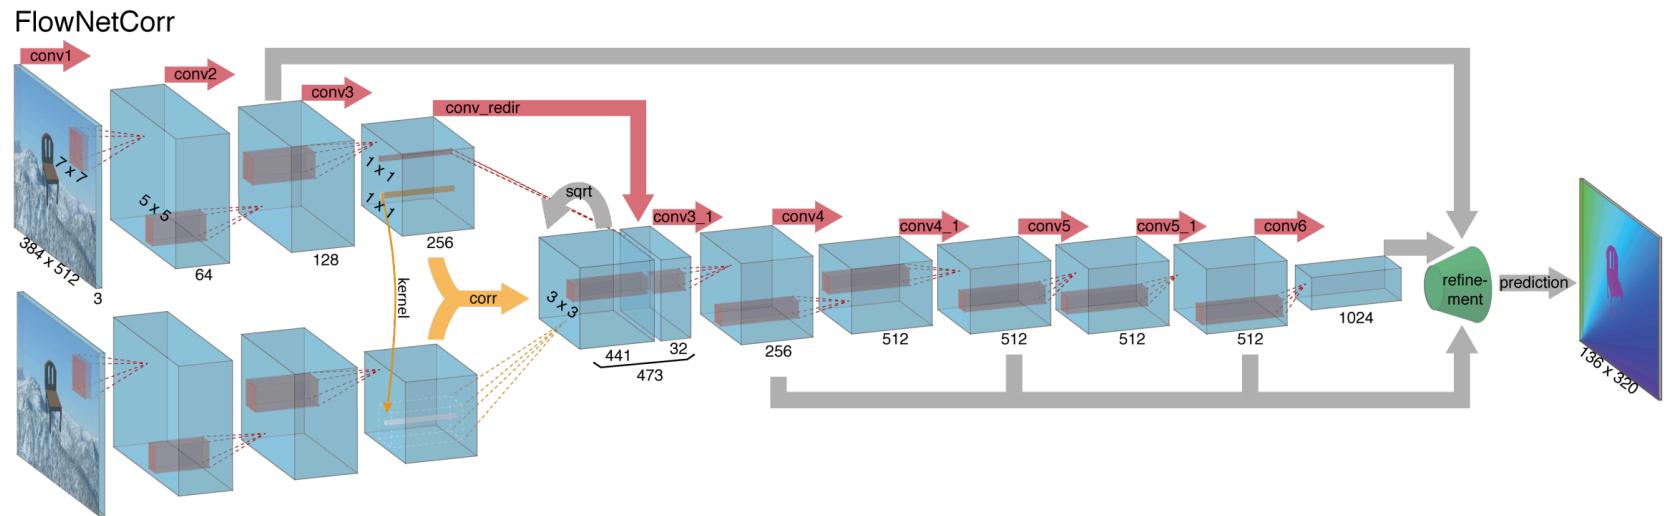
[FlowNet: Learning Optical Flow with Convolutional Networks](#)

FlowNetSimple



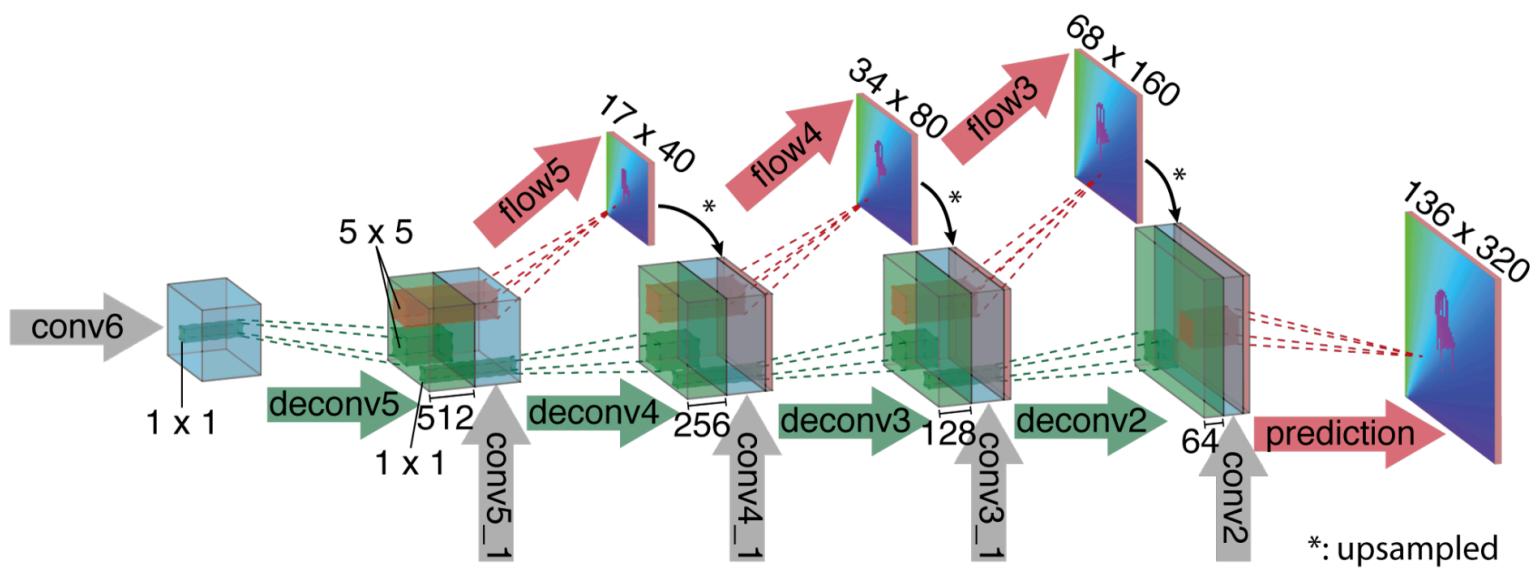
FlowNet: Learning Optical Flow with Convolutional Networks

FlowNetCorr



FlowNet: Learning Optical Flow with Convolutional Networks

Refinement of the coarse feature maps



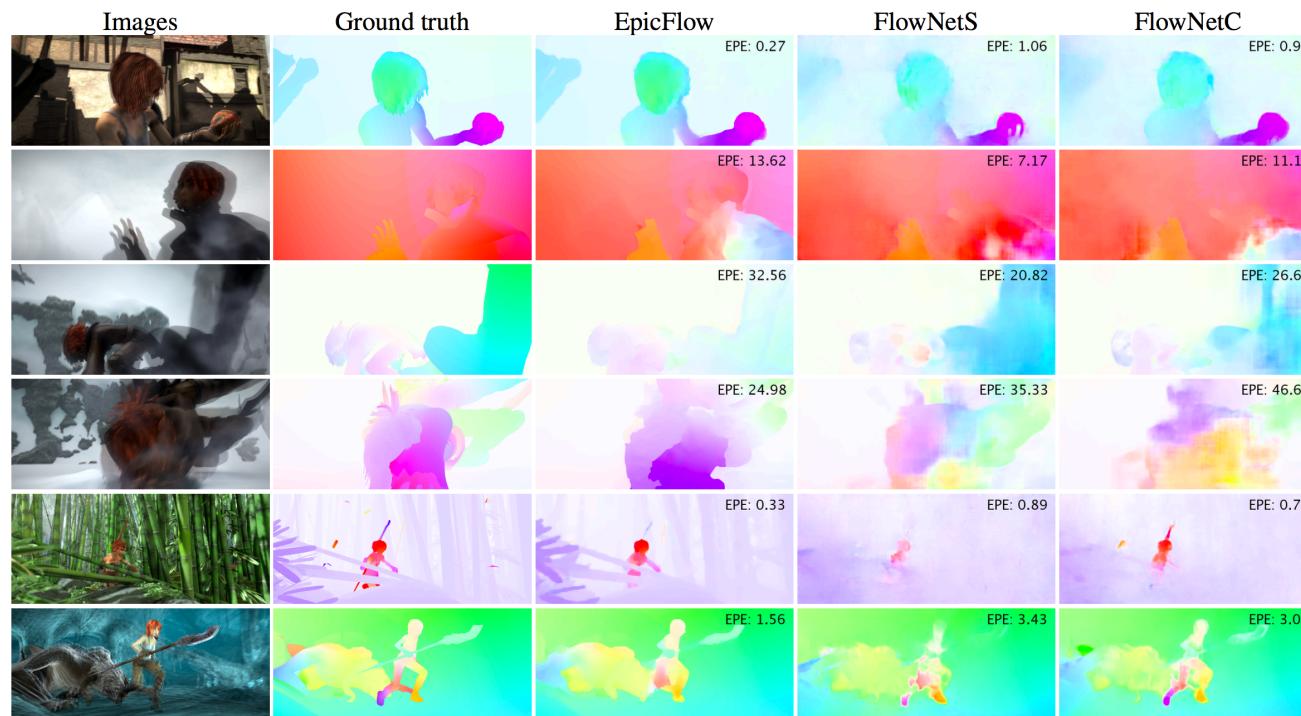
FlowNet: Learning Optical Flow with Convolutional Networks

Examples of Data Pairs



[FlowNet: Learning Optical Flow with Convolutional Networks](#)

Results



FlowNet: Learning Optical Flow with Convolutional Networks